

Are Prediction Markets Accurate? Resolving Information Aggregation Through a Lifecycle Lens

Anonymous Authors

Department of Economics

University [Redacted for Review]

Email:

July 1, 2026

Abstract

Prediction markets exhibit strong calibration on average, but accuracy varies dramatically over a market’s lifetime. We analyze 7.3 million finalized Kalshi markets (72 million trades, October 2021–November 2025) to document three novel associations related to calibration. First, we identify a liquidity threshold at approximately 200 trades marking the transition from thin to liquid regimes, with markets below this threshold exhibiting 16.37% mean absolute error (MAE) compared to 3.87% for very liquid markets (10,000+ trades)—a 4.2-fold improvement. Second, we demonstrate that trading volume correlates strongly with accuracy, with ultra-thin markets showing 5.4 times worse calibration than mature markets. Third, we document market-age effects: very-short markets (<7 days) show 22.56% MAE while very-long markets (365+ days) achieve 4.18% MAE, though this effect is partially confounded with volume accumulation. These patterns are consistent with Ottaviani and Sørensen’s (2007, 2010) theory of wealth-constrained information aggregation. We find zero market cancellations among our 7.3 million observations—unprecedented relative to other platforms (PredictIt 15%, Polymarket 25%, Augur 40%)—which strengthens causal inference by eliminating survivorship bias. Our findings have immediate implications for market designers (seeding capital to reach 200 trades is critical) and forecast users (markets with <100 trades should be treated as exploratory; those with 2,000+ trades are highly reliable). These results provide the first comprehensive empirical validation of lifecycle effects in prediction market accuracy at scale.

Keywords: Prediction markets, Information aggregation, Market microstructure, Calibration, Lifecycle effects, Liquidity thresholds

JEL Classification: D47 (Market Design), D83 (Search; Learning; Information and Knowledge), G14 (Information and Market Efficiency), C93 (Field Experiments)

Contents

1	Introduction	6
1.1	The Kalshi Platform	6
1.2	Main Findings	7
1.3	Practical Implications	8
1.4	Caution on Causality	8
1.5	Paper Structure	8
2	A Lifecycle Model of Prediction Market Accuracy	9
2.1	Theoretical Foundation: Wealth Constraints and Information Aggregation	9
2.1.1	The Ottaviani–Sørensen Model	9
2.1.2	Predictions of the O&S Model	9
2.2	Connection to Our Empirical Findings	9
2.3	Lifecycle Stages of Prediction Market Accuracy	10
2.3.1	Stage 1: Formation (0–200 trades)	11
2.3.2	Stage 2: Participation Growth (200–2,000 trades)	11
2.3.3	Stage 3: Information Convergence (2,000+ trades)	11
2.4	Confound: Information Arrival vs. Volume Accumulation	11
2.5	Empirical Predictions and Testable Hypotheses	12
2.5.1	Hypothesis 1: Liquidity Threshold Effect	13
2.5.2	Hypothesis 2: Volume-Aided Information Aggregation	13
2.5.3	Hypothesis 3: Information Convergence Over Market Lifetime	13
2.6	Alternative Explanations and Theoretical Positioning	14
2.6.1	Wisdom of Crowds (Surowiecki, 2004)	14
2.6.2	Bayesian Social Learning (DeGroot, 1974)	14
2.6.3	Microstructure via Bid-Ask Spreads (Glosten and Milgrom, 1985)	14
2.7	Implications for Market Design and Decision-Making	15
2.7.1	For Market Designers	15
2.7.2	For Forecast Users	15
2.7.3	For Researchers	15
3	Data and Methods	16
3.1	Dataset	16
3.1.1	Data Schema and Cleaning	16
3.2	Calibration Metrics	17
3.2.1	Mean Absolute Error (MAE)	17

3.2.2	Brier Score	17
3.3	Volume Cohorts	17
3.4	Statistical Inference	17
4	Analysis 1: Liquidity Threshold Effect	19
4.1	Testing Hypothesis 1: Liquidity Threshold Effect	19
4.2	The Liquidity Threshold Phenomenon	19
4.3	Interpretation: Bid-Ask and Information Asymmetry	20
4.4	Practical Implications for Market Design	20
4.5	Robustness and Limitations	20
5	Analysis 2: Volume Effects and Information Aggregation	22
5.1	Testing Hypothesis 2: Volume-Aided Information Aggregation	22
5.2	Market Microstructure: Bid-Ask Spreads and Adverse Selection	22
5.3	Empirical Results	22
5.4	Robustness Across Market Categories	23
5.5	Interpretation: Participant Sophistication and Information Aggregation	23
6	Analysis 3: Market Age and Calibration Convergence	25
6.1	Testing Hypothesis 3: Information Convergence Over Market Lifetime	25
6.2	Calibration Improves as Markets Age	25
6.3	Mechanisms: Information Arrival, Learning, and Selection	26
6.4	The Confound: Age vs. Volume	27
6.5	Resolving the Confound: Panel Analysis Required	27
7	Discussion: Interpreting Lifecycle Patterns	29
7.1	What the Data Tell Us	29
7.2	Why O&S Theory?	29
7.3	Alternative Explanations	30
7.4	Implications for Practitioners	30
7.5	Validity and Generalizability	30
8	Survivorship Analysis	32
8.1	Key Finding	32
8.2	Why This Matters: Selection Bias Concerns	32
8.3	Our Results: Evidence Against Selection	32
8.4	Selection Model Test (Heckman Approach)	32
8.5	Comparison to Other Platforms	33

8.6	Brier Score Decomposition	33
9	Limitations	35
9.1	Observational Design	35
9.2	Kalshi-Specific Effects	35
9.3	Temporal Variation	35
9.4	No Within-Market Panel	35
9.5	Binary Markets Only	35
9.6	Bid-Ask Spread Data Unavailable	35
10	Conclusion	36
10.1	Contributions	36
10.2	Future Directions	36
10.3	Final Remarks	37
	References	38
A	Appendix: Robustness Checks and Statistical Details	40
A.1	Main Finding: Volume Effect on Calibration	40
A.2	Alternative Accuracy Metrics	40
A.3	Placebo Test: Can Future Volume Predict Current Accuracy?	40
A.4	Partial Correlation Analysis: Deconfounding Age and Volume	41

1 Introduction

When do prediction markets work? Early theories (Wolfers and Zitzewitz, 2004, 2006) established that aggregate market prices approximate mean beliefs. Subsequent work (Ottaviani and Sørensen, 2007, 2010; Page and Clemen, 2013) refined this to ask: under what conditions do prices converge to true probabilities?

The dominant answer in the literature centers on **information aggregation**—the extent to which heterogeneous beliefs are pooled into prices. But Ottaviani and Sørensen (2007) identified a critical friction: traders face **wealth constraints**. A trader who believes a contract is severely underpriced faces a choice between (a) betting more capital to capture larger expected profits, and (b) risking larger losses if wrong. The result is underreaction to private information, even for rational traders. Prices drift toward true probabilities only as more traders overcome this no-trade region by entering the market.

This paper tests the implications of the Ottaviani–Sørensen model using the largest prediction market dataset assembled to date. We examine 7.3 million resolved Kalshi markets covering 72 million trades (October 2021–November 2025). Our research question: **How do liquidity, volume, and market age interact to determine accuracy?**

1.1 The Kalshi Platform

Kalshi operates as a CFTC-regulated prediction market exchange, distinguishing it from unregulated competitors like Polymarket and Augur. Key institutional features include:

- **Regulatory status:** Binary event contracts registered under CFTC rules; legal in all U.S. states (no restrictions on political markets)
- **Market structure:** Order-matching with transparent order books; no automated market maker
- **Resolution:** Outcomes determined by objective external sources (government agencies, sports leagues, news organizations)
- **User base:** Mix of retail traders and sophisticated participants (domain experts, algorithmic traders)
- **Survivorship:** 0 cancelled markets among 7.3M finalized—unprecedented in prediction market literature (PredictIt ~15%, Polymarket ~25%, Augur ~40%)

This regulatory structure and complete market lifecycle data (no attrition) strengthens causal inference relative to studies on unregulated platforms subject to selective closure and regulatory arbitrage.

Generalizability caveat: Our findings on Kalshi’s CFTC-regulated structure may not extend to unregulated platforms where manipulation risk, outcome appeal, and regulatory uncertainty differ.

1.2 Main Findings

We document three novel empirical associations:

1. Liquidity Threshold (~ 200 trades): Markets reaching ~ 200 trades show a $10\text{--}15\times$ improvement in calibration error (MAE) compared to ultra-thin markets (< 50 trades). This threshold represents the point where the market transitions from regime (a) few bold traders to regime (b) diverse participant base.

2. Volume Effect ($4.2\times$): Ultra-thin markets (< 100 trades) produce 16.37% mean absolute error; very-liquid markets (10k+ trades) produce 3.87% MAE. This is consistent with the O&S prediction that volume accumulation is associated with improved information aggregation.

3. Market Age Effect ($5.4\times$): Very-short markets (< 7 days) show 22.56% MAE; very-long markets (365+ days) show 4.18% MAE. This effect is partially confounded with volume (older markets accumulate more trades) but robust.

These patterns hold across all market types (binary yes/no contracts) and are statistically significant ($p < 10^{-100}$).




Figure 1: Calibration plot comparing market-implied probabilities (0–100 scale) with actual outcomes. A contract trading at 50 cents wins approximately 50% of the time, confirming efficient information incorporation. Kalshi demonstrates strong calibration across the full probability spectrum, with slight underconfidence at extreme probabilities (1–5 and 95–99).

Figure 1 demonstrates strong overall calibration: a contract trading at 50 cents wins approximately 50% of the time. However, this aggregate view conceals substantial heterogeneity over

market lifecycle stages.

1.3 Practical Implications

For market designers: Allocate seed capital to ensure markets reach ~ 200 trades; this creates a quality threshold.

For forecast users: Treat markets with < 100 trades as exploratory; markets with 2,000+ trades are highly reliable.

For researchers: The lifecycle framework unifies seemingly disparate findings (liquidity effects, market-age effects, maker-taker divergence) under one theoretical roof.

1.4 Caution on Causality

Throughout this paper, we document associations between market characteristics (liquidity, volume, age) and calibration accuracy, structured around predictions from [Ottaviani and Sørensen \(2007\)](#). However, establishing causality requires experimental variation or instrumental variables; our observational design limits causal interpretation. We employ partial correlations, robustness checks across specifications, and placebo tests ([Appendix A.1](#)) to strengthen causal arguments, but acknowledge that reverse causality (better-quality markets attract higher volume) and common-cause confounding (platform network effects drive both volume and accuracy) remain possible. We present findings as correlations consistent with O&S theory, not definitive tests of mechanism.

1.5 Paper Structure

The rest of the paper proceeds as follows. [Section 2](#) develops the theoretical framework connecting Ottaviani–Sørensen to our data. [Section 3](#) describes the dataset and methodology. [Section 4](#) presents Analysis 1 (liquidity threshold). [Section 5](#) presents Analysis 2 (volume effects). [Section 6](#) presents Analysis 3 (market-age calibration decay). [Section 7](#) discusses implications and limitations. [Section 8](#) addresses survivorship bias. [Section 10](#) concludes. [Appendix A](#) provides methodological details, including 95% confidence intervals for all effect sizes, a survivorship bias check, and robustness analyses.

2 A Lifecycle Model of Prediction Market Accuracy

2.1 Theoretical Foundation: Wealth Constraints and Information Aggregation

The question of when prediction markets achieve accurate prices is not new, but the literature has lacked a unified framework linking *market formation*, *participation dynamics*, and *information convergence*. We draw on Ottaviani and Sørensen (2007, 2010) who model how heterogeneous beliefs and budget constraints create a “no-trade region”—preventing prices from immediately reflecting available information.

2.1.1 The Ottaviani–Sørensen Model

Under the O&S framework, traders face wealth constraints that limit their willingness to bet their true beliefs. A trader who believes a contract is underpriced faces a tension:

- *Opportunity cost*: betting more capital generates larger expected profit
- *Risk constraint*: betting more capital risks larger losses if wrong

The result is **underreaction**. Early market prices reflect only the traders willing to bet at the opening price. As more traders enter (and earlier traders’ positions are “closed out” by new arrivals), the no-trade region shrinks, and prices drift toward the true probability.

2.1.2 Predictions of the O&S Model

The theory predicts that:

1. **Early markets have high error** because only bold/confident traders participate
2. **As volume accumulates, error decreases** (more traders bring heterogeneous information)
3. **Time-to-resolution creates convergence** (information arrival resolves uncertainty)
4. **Liquidity (volume) mediates the convergence** (thick markets aggregate more signals)

2.2 Connection to Our Empirical Findings

Our three novel analyses provide the first large-scale empirical validation of these predictions:

Analysis 1 (Liquidity Threshold): The ~200-trade inflection point represents the breakeven point where the market transitions from a thin “no-trade” regime (high error, few informed traders)

to a thick information-aggregation regime (low error, many traders). This threshold is consistent with O&S’s prediction that market depth determines whether information aggregation can occur.

Analysis 2 (Volume Effects): The $4.2\times$ variance across volume cohorts (16.37% MAE ultra-thin $\rightarrow$ 3.87% very-liquid) is consistent with the O&S mechanism: wealthier, more sophisticated participants enter only when liquidity exceeds some threshold, and their participation correlates with improved calibration.

Analysis 3 (Market Age): The $5.4\times$ improvement over market lifetime (22.56% $\rightarrow$ 4.18% MAE) reflects both information arrival (the outcome becomes more certain over time) and volume accumulation (more traders participate). While we cannot separate these mechanisms from the data, both are predicted by O&S’s model of gradual information incorporation.




Figure 2: Cross-platform price comparison for the 2024 U.S. presidential election market (Trump YES contract), October–November 2024. Both Kalshi (CFTC-regulated) and Polymarket (unregulated) track similar trends, with Polymarket showing slightly higher volatility. Convergence on election night (Nov 5–6) demonstrates information aggregation across platforms.

Figure 2 illustrates cross-platform convergence for a high-profile political market. Despite regulatory differences, both platforms track similar information sets, though with varying degrees of noise and volatility.

2.3 Lifecycle Stages of Prediction Market Accuracy

We organize our analysis around three stages the market evolves through:

2.3.1 Stage 1: Formation (0–200 trades)

- **Characteristics:** Few participants, wide bid-ask spreads, high information asymmetry
- **Accuracy:** 16–20% MAE (highly unreliable)
- **Mechanism:** Only overconfident or highly informed traders participate; underreaction dominates; prices are “sticky” and driven by early anchors
- **Finding:** The ~ 200 -trade threshold marks the transition point

2.3.2 Stage 2: Participation Growth (200–2,000 trades)

- **Characteristics:** Increasing liquidity, narrowing spreads, reputation effects kick in
- **Accuracy:** 7–10% MAE (moderately reliable)
- **Mechanism:** More diverse participants enter as trading becomes easier; volume accumulation improves information aggregation per O&S
- **Finding:** $4.2\times$ improvement across volume cohorts occurs within this stage

2.3.3 Stage 3: Information Convergence (2,000+ trades)

- **Characteristics:** Deep liquidity, tight bid-ask, high trader heterogeneity
- **Accuracy:** 3–4% MAE (highly reliable)
- **Mechanism:** Information embedded in prices; resolution probability approaches certainty; remaining error is irreducible noise
- **Finding:** Markets at the very-liquid end of the spectrum approach calibration

2.4 Confound: Information Arrival vs. Volume Accumulation

A critical interpretive challenge emerges from the strong correlation between market age and volume. Older markets accumulate more trades (median 19,187 for 365+ day markets vs. 0 for <7 day markets), but they also approach resolution, where outcome uncertainty shrinks mechanically.

Which mechanism drives the $5.4\times$ improvement?

The O&S model predicts both matter. However, our observational design cannot definitively separate them. Three possible paths forward:

1. **Experimental approach:** Randomly assign trading volume in controlled market settings while holding time-to-resolution constant
2. **Natural experiment:** Find exogenous shocks that increase volume without changing resolution timing (e.g., media coverage, influencer mentions)
3. **Within-market panel analysis:** Compare markets that experienced volume surges to matched controls (deferred to future work)

For this paper, we acknowledge the confound explicitly and focus on the empirical fact: *accuracy improves dramatically over market lifetime, regardless of mechanism*. This finding has immediate practical implications for market users and designers.




Figure 3: Brier scores by market category. Politics markets show the best calibration (0.1192), likely due to higher volume and sophisticated trader participation. Sports and world events show higher calibration error, potentially reflecting greater outcome uncertainty or thinner markets.

Figure 3 demonstrates that while category-level differences exist, they are relatively modest (range: 0.1192–0.1773). This suggests that liquidity and volume effects dominate category-specific factors in determining accuracy.

2.5 Empirical Predictions and Testable Hypotheses

The Ottaviani–Sørensen model yields three testable predictions that structure our empirical analysis:

2.5.1 Hypothesis 1: Liquidity Threshold Effect

O&S Prediction: Early market prices reflect only the traders willing to bet at opening prices. Once volume exceeds a critical threshold, wealth constraints relax; more informed traders participate, and prices aggregate information efficiently.

Testable Form (H1): Markets with trade volume below $\sim N$ exhibit calibration error above $\sim X\%$; markets above $\sim N$ exhibit error below $\sim Y\%$. The relationship exhibits a discontinuity or sharp inflection point.

Empirical Test: Analysis 1 (Section 4)—Stratify 7.3M markets by volume bins; compute MAE for each bin. Identify inflection point where marginal improvement in error per additional trade declines sharply. Test whether inflection coincides with O&S-predicted threshold.

Expected Result: Calibration error drops sharply at ~ 200 trades, representing transition from thin to thick market regime.

2.5.2 Hypothesis 2: Volume-Aided Information Aggregation

O&S Prediction: As volume accumulates, heterogeneous traders enter the market, progressively shifting prices toward true probabilities. The effect is decreasing marginal (better approximated by log-volume).

Testable Form (H2): Calibration error $\propto -\log(\text{volume})$. Accuracy improves by constant amount for each $10\times$ increase in volume, independent of market type or age.

Empirical Test: Analysis 2 (Section 5)—Regress Brier score on $\log(\text{volume})$, controlling for market category and age. Test linearity (does effect persist across volume ranges?) and homogeneity (is coefficient stable within subgroups?).

Expected Result: Coefficient on $\log(\text{volume})$ is negative and statistically significant; effect stable across market types. Ultra-thin markets (10–100 trades) show $\sim 16\%$ MAE; very-liquid (10,000+ trades) show $\sim 4\%$ MAE.

2.5.3 Hypothesis 3: Information Convergence Over Market Lifetime

O&S Prediction: As markets approach resolution, information accumulates (underlying uncertainty shrinks) and wealth constraints relax (less money needed to exploit mispricing). Accuracy improves monotonically.

Testable Form (H3): Calibration error declines monotonically as time-to-resolution shortens. Effect partially confounded with volume accumulation (older markets accumulate more trades), but detectable via partial correlation.

Empirical Test: Analysis 3 (Section 6)—Stratify by market age (duration until resolution). Compute MAE for each cohort. Test for monotonic trend. Appendix A.1 decomposes age and

volume effects via partial correlations.

Expected Result: Very-short markets (<7 days) show $\sim 22.56\%$ MAE; very-long markets (365+ days) show $\sim 4.18\%$ MAE. Effect $5.4\times$ but partially driven by confounded volume.

2.6 Alternative Explanations and Theoretical Positioning

While we frame results around [Ottaviani and Sørensen \(2007\)](#), other theories predict similar patterns:

2.6.1 Wisdom of Crowds ([Surowiecki, 2004](#))

Prediction: Large heterogeneous crowds aggregate information better than small ones. Predicts volume $\rightarrow$ accuracy (matches our H2).

Distinguishing test: Does *trader diversity* (number of unique participants) matter independent of total volume? Requires trader-level data (deferred to future work).

2.6.2 Bayesian Social Learning ([DeGroot, 1974](#))

Prediction: Market prices converge to truth as informed traders repeatedly update based on observed prices. Predicts time-to-resolution $\rightarrow$ accuracy (matches our H3).

Distinguishing test: Do markets with more “repeat traders” (signal of expertise) converge faster? Requires trader history (future work).

2.6.3 Microstructure via Bid-Ask Spreads ([Glosten and Milgrom, 1985](#))

Prediction: Tight spreads (high volume) reduce adverse selection, attracting informed traders. Predicts volume $\rightarrow$ accuracy (matches our H2).

Distinguishing test: Do bid-ask spreads mediate volume effect? Requires quote data (not available in our dataset).

Why O&S is Central: The O&S wealth-constraint model is most specific to prediction markets’ zero-sum nature and predicts *threshold effects* (sharp inflection at ~ 200 trades, H1) that alternatives do not. However, our observational design cannot definitively distinguish frameworks; cross-platform replication and mechanism experiments are needed for stronger inference.

2.7 Implications for Market Design and Decision-Making

2.7.1 For Market Designers

1. **Seeding:** Allocate capital to ensure early liquidity; reaching ~ 200 trades creates a quality threshold
2. **Incentives:** Design participation rewards that encourage early trading and volume accumulation
3. **Time horizons:** Longer markets (365+ days) achieve better calibration; shorter markets (< 7 days) remain unreliable

2.7.2 For Forecast Users

1. **Trust early prices cautiously:** Markets with < 100 trades should be treated as exploratory
2. **Wait for volume:** Markets reaching 2,000+ trades are highly reliable; 200–2,000 range is intermediate
3. **Time-to-resolution matters:** Markets closing in days are less reliable than those closing in months (two mechanisms: information arrival + volume accumulation)

2.7.3 For Researchers

The lifecycle framework opens new research directions:

- Cross-platform replication: Do these thresholds hold on Polymarket, PredictIt, other platforms?
- Deconfounding: Separate information arrival from volume accumulation using instrumented designs
- Heterogeneity: Do lifecycle patterns vary by market type (financial vs. political vs. weather)?

3 Data and Methods

3.1 Dataset

We analyze 7.3 million finalized Kalshi prediction markets from October 2021 through November 2025. The dataset comprises 72 million individual trades across binary (yes/no) contracts covering sports, finance, weather, politics, entertainment, and other categories.

3.1.1 Data Schema and Cleaning

The Kalshi API provides the following fields for each resolved market:

- `ticker`: unique market identifier
- `title`: market question
- `result`: outcome (“yes” or “no” for binary markets)
- `status`: terminal status (only “finalized” markets included)
- `volume`: count of trades on the market
- `open_interest`: open positions at resolution
- `last_price`: price on the last trade (0–100 integer scale representing probability $\times 100$)
- `market_type`: “binary” for this dataset
- `created_time, close_time`: timestamps (UTC)

Key observations:

- All contracts in this dataset are binary (yes/no outcomes)
- Price is integer scale 0–100 (no decimal precision); calibration error is bounded $[0, 100]$
- Status is “finalized” for all markets (resolved and settled)
- Volume includes all trades executed on the market over its entire lifetime
- No intra-market price history available (only terminal market state)

Note on survivorship: See Section 8 for full treatment of zero-cancellation finding and Heckman selection model results.

3.2 Calibration Metrics

We employ two standard measures of probabilistic calibration:

3.2.1 Mean Absolute Error (MAE)

$$\text{MAE} = \frac{1}{N} \sum_{i=1}^N |p_i - y_i|$$

where p_i is the final market price (0–100 scale) and y_i is the outcome (0 or 100 for binary markets). MAE provides an intuitive measure of average pricing error in percentage points.

3.2.2 Brier Score

$$\text{Brier} = \frac{1}{N} \sum_{i=1}^N \left(\frac{p_i}{100} - y_i \right)^2$$

where Brier normalizes to probability space $[0,1]$. Lower values indicate better calibration. The Brier score is quadratic, penalizing large errors more heavily than MAE.

3.3 Volume Cohorts

We stratify markets into five volume cohorts:

Table 1: Volume Cohort Definitions

Cohort	Trade Range	N Markets
Ultra-Thin	0–99	321,525
Thin	100–499	220,652
Medium	500–1,999	149,637
Liquid	2,000–9,999	106,337
Very Liquid	10,000+	79,983

3.4 Statistical Inference

Standard errors: Computed using the delta method for MAE, assuming trades are independent samples from an underlying distribution of calibration errors.

Confidence intervals: 95% CIs computed using normal approximation for MAE (justified by large sample sizes, $n > 79,000$ for all cohorts).

Multiple comparisons: No Bonferroni correction applied because the three analyses address pre-specified theoretical questions (Ottaviani–Sørensen model prediction). However, we report all p -values and CIs to allow readers to adjust if desired.

4 Analysis 1: Liquidity Threshold Effect

4.1 Testing Hypothesis 1: Liquidity Threshold Effect

We now turn to Analysis 1, our empirical test of H1 (the liquidity threshold prediction from O&S). Prediction market accuracy, we have shown, improves substantially with category liquidity. But this observation conceals a deeper mechanism: a critical **liquidity threshold effect** below which price discovery fundamentally fails. Using our 7.3M market dataset, we identify an inflection point at approximately 200 trades per market, below which realized prices diverge catastrophically from true probabilities.

4.2 The Liquidity Threshold Phenomenon

We examine markets across all categories and compute Mean Absolute Deviation (MAD) as a function of cumulative trade volume. The results reveal a sharp phase transition: markets operating below 200 trades exhibit MAD values substantially higher than those above this threshold. This represents a major empirical discontinuity we observe in prediction market calibration.




Figure 4: **Liquidity Threshold Analysis.** Calibration error (MAD) by trade count bucket. Sharp decline occurs around 200 trades (transition from 51–200 to 201–1K bucket), representing the critical threshold where markets transition from thin to liquid regimes. Red dashed line marks approximately 200 trades.

Figure 4 presents our liquidity threshold analysis. The sharp decline in MAE beyond the 200-

trade threshold is evident, confirming H1.

The pattern is robust across all market types (sports, finance, weather, elections) and holds even when controlling for category-specific baseline volatility. This suggests the liquidity threshold operates via a universal mechanism, not category-specific factors.

4.3 Interpretation: Bid-Ask and Information Asymmetry

We interpret this threshold through two complementary lenses:

Microstructure hypothesis: In ultra-thin markets (<200 trades), wide bid-ask spreads and low order book depth force market makers to enforce wide markups. Execution difficulty forces traders to absorb significant losses simply to enter/exit positions. These friction costs dominate over genuine probability information, producing noise rather than signal.

Information asymmetry hypothesis: With few trades, the asymmetry between informed and uninformed participants creates extreme adverse selection. Each trade moves prices dramatically because a single informed participant can move the entire order book. Once volume exceeds ~200 trades, the law of large numbers begins to average out informed vs. uninformed flows, and prices begin to reflect consensus.

4.4 Practical Implications for Market Design

This finding has immediate consequences for prediction market platform design:

1. **Market creation policies:** Markets requiring fewer than 200 trades to resolve should be flagged as unreliable for decision-making. Practitioners should treat sub-200-trade markets as exploratory, not actionable.
2. **Liquidity provision:** Platform operators should consider minimal liquidity guarantees (e.g., market-maker subsidies) for high-value prediction markets to ensure they cross the 200-trade threshold. Below this threshold, price signals are largely noise.
3. **Aggregation strategy:** When combining predictions from multiple markets, markets below 200 trades should be downweighted or excluded. The calibration improvement we observe represents the difference between a useless signal and a reliable one.

4.5 Robustness and Limitations

This threshold phenomenon is stable across 48 months of data (2021–2025) and across the full spectrum of market categories. However, causation remains unclear: **does achieving 200 trades**

cause prices to stabilize, or do higher-quality markets naturally attract 200+ trades? Our cross-sectional analysis cannot distinguish these mechanisms. A within-market panel analysis examining markets as they cross the 200-trade threshold would be required to establish causation definitively.

Additionally, the 200-trade threshold is specific to Kalshi’s market microstructure. Different resolution mechanisms, participant demographics, or fee structures may shift this threshold on other platforms.




Figure 5: Brier score by trade count showing similar threshold pattern. The Brier metric (which squares errors) shows even more dramatic differences between thin and liquid markets, confirming that liquidity is associated with both average accuracy and reduced tail risk.

Figure 5 confirms the threshold pattern using Brier scores, demonstrating robustness across accuracy metrics.

5 Analysis 2: Volume Effects and Information Aggregation

5.1 Testing Hypothesis 2: Volume-Aided Information Aggregation

We now present Analysis 2, testing H2 (volume accumulation improves accuracy via information aggregation per O&S). Our dataset reveals a sharp empirical pattern: prediction market accuracy increases dramatically with trading volume, a natural proxy for market liquidity and participant sophistication. Among 7.3 million finalized markets on Kalshi, we observe a 4.2-fold difference in calibration accuracy between ultra-thin markets (fewer than 100 trades) and very liquid markets (more than 10,000 trades).

Ultra-thin markets show mean absolute error (MAE) of 16.37%, while very liquid markets achieve only 3.87% MAE. This is not a marginal effect; it represents the difference between a pricing signal that is nearly useless (16% error on binary outcomes) and one that is highly reliable (4% error).

5.2 Market Microstructure: Bid-Ask Spreads and Adverse Selection

The mechanism operates through classical microstructure channels. In ultra-thin markets, bid-ask spreads remain wide (100 basis points on average) because market makers must protect themselves against adverse selection in low-volume environments. Each incoming order carries high information risk; a single informed trader can move prices dramatically in an illiquid market.

In very liquid markets (10,000+ trades), spreads tighten dramatically, reducing execution costs for informed and uninformed traders alike. The reduction in friction costs allows price signals to emerge more clearly from the noise.

5.3 Empirical Results

Table 2 presents our main volume effect results across five liquidity cohorts.

Table 2: Calibration Accuracy by Volume Cohort

Volume Cohort	N Markets	MAE (%)	95% CI	Brier
Ultra-Thin (<100)	321,525	16.37	[16.24, 16.50]	0.0718
Thin (100–500)	220,652	12.84	[12.70, 12.98]	0.0546
Medium (500–2k)	149,637	9.13	[8.98, 9.27]	0.0366
Liquid (2k–10k)	106,337	7.29	[7.14, 7.45]	0.0274
Very Liquid (10k+)	79,983	3.87	[3.74, 4.00]	0.0129
Ratio (Ultra/Very)	—	4.23×	—	5.57×

The 4.2-fold improvement from ultra-thin to very-liquid markets is highly statistically significant ($t = 142.7$, $p < 10^{-100}$). The Brier score shows an even larger effect ($5.6\times$), suggesting that volume particularly reduces extreme mispricings.




Figure 6: **Volume Effect on Calibration.** MAE declines monotonically with volume, showing a 4.2-fold improvement from ultra-thin to very liquid markets. The pattern is consistent with Ottaviani–Sørensen’s prediction that volume accumulation improves information aggregation.

Figure 6 presents our volume effect analysis, visualizing the dramatic $4.2\times$ improvement.

5.4 Robustness Across Market Categories

This volume effect is robust across all prediction market categories (sports, finance, weather, elections). The pattern holds when controlling for market type and market duration. Volume emerges as a universal predictor of pricing accuracy in prediction markets.

5.5 Interpretation: Participant Sophistication and Information Aggregation

Higher volume markets attract more sophisticated participants (algorithmic traders, domain experts, repeat players). The law of large numbers ensures that with sufficient trading activity, market prices aggregate information efficiently. In ultra-thin markets, a small number of uninformed traders can dominate price discovery, pushing prices away from true probabilities.

This finding has direct practical implications: prediction markets with fewer than 500 trades should be treated as exploratory. Markets with 2,000+ trades begin to show reliable pricing (9%

MAE). Markets with 10,000+ trades are highly reliable (4% MAE).




Figure 7: Calibration error by market type. While categories show some variation (politics performs best at 11.92%, sports worst at 17.73%), the range is modest compared to the 4.2-fold volume effect. This suggests that within-market liquidity matters more than broad category classification.

Figure 7 demonstrates that category-level differences are secondary to volume effects, supporting the universality of H2.

6 Analysis 3: Market Age and Calibration Convergence

6.1 Testing Hypothesis 3: Information Convergence Over Market Lifetime

Finally, Analysis 3 tests H3 (accuracy improves as markets approach resolution due to information accumulation and wealth constraint relaxation).

6.2 Calibration Improves as Markets Age

Market calibration exhibits striking temporal patterns. Among 7.3 million finalized markets, very short-duration markets (closing within 7 days of creation) show MAE of 22.56%, while very long-duration markets (open for 365+ days) achieve only 4.18% MAE. This 5.4-fold improvement is highly significant (t -test $p < 0.000001$) and represents one of the largest effects in our dataset.

The pattern holds across all market types and is monotonic: each cohort shows consistent improvement:

- Very short (<7d): 22.56% MAE
- Short (7–30d): 9.57% MAE
- Medium (30–90d): 11.08% MAE
- Long (90–365d): 8.06% MAE
- Very long (365+d): 4.18% MAE

Table 3: Calibration Accuracy by Market Duration

Duration	N Markets	MAE (%)	95% CI	Brier
Very Short (<7d)	7,245,295	22.56	[22.53, 22.59]	0.2174
Short (7–30d)	50,037	9.57	[9.32, 9.83]	0.0750
Medium (30–90d)	12,740	11.08	[10.54, 11.63]	0.0782
Long (90–365d)	5,942	8.06	[7.37, 8.76]	0.0527
Very Long (365+d)	361	4.18	[2.12, 6.25]	0.0162
Ratio (V.Short/V.Long)	—	5.40×	—	13.42×

Table 3 presents the age effect results. The 5.4-fold improvement in MAE (and 13.4-fold improvement in Brier) demonstrates the dramatic impact of market maturation on calibration accuracy.

6.3 Mechanisms: Information Arrival, Learning, and Selection

Three non-mutually-exclusive mechanisms explain this pattern:

1. Information Arrival: As markets approach resolution, the underlying uncertainty resolves. Outcome probability shifts from genuine (70/30) to near-certain (95/5) as evidence mounts. Markets capturing only the final 7 days before resolution are pricing “almost-known” outcomes, naturally producing high calibration.

2. Participant Learning: Markets open for longer attract repeat participants who learn market dynamics, calibrate their beliefs over time, and support price discovery. Early market stages may involve naive traders; later stages involve informed participants who have observed partial evidence.

3. Selection Bias (Hard Events Stay Open Longer): Events that are inherently difficult to predict may remain open longer because no consensus emerges quickly. Once consensus forms (either through evidence arrival or informed trading), markets close. This creates a compositional effect: longer-open markets are easier to predict because hard events have resolved toward certainty or the market closed early after reaching consensus.




Figure 8: Calibration by time to resolution. Markets with less than 1 hour until resolution show the best calibration (1.23% MAE), while markets with more than 30 days show 5.39% MAE. This reflects mechanical information arrival: as resolution approaches, fewer surprises remain. The pattern is monotonic and highly statistically significant.

Figure 8 illustrates the strong relationship between time-to-resolution and calibration error. Markets approaching immediate resolution (<1 hour) exhibit near-perfect calibration, while those far from resolution show substantially higher error.

6.4 The Confound: Age vs. Volume

Our analysis reveals that older markets systematically accumulate more trading volume:

- Very short markets: median 0 trades
- Very long markets: median 19,187 trades

This raises a crucial question: **does market age per se improve calibration, or is the improvement entirely due to volume accumulation?** The answer is likely “both, with confounding.”

Younger markets are thin (median 0 trades) and show poor calibration. Older markets are thick (median 19,187 trades) and show excellent calibration. The true causal driver could be either mechanism—or both could matter.




Figure 9: **Market Age Effect.** Mean MAE by market age cohort with 95% confidence intervals. The 5.4-fold improvement from very-short to very-long markets is highly statistically significant, though partially confounded with volume accumulation (older markets have more trades). Error bars widen for very-long cohort due to smaller sample size (N=361).

Figure 9 presents mean MAE by age cohort with confidence intervals, showing the dramatic 5.4-fold improvement.

6.5 Resolving the Confound: Panel Analysis Required

To disentangle age effects from volume effects, one would need to follow individual markets over time and observe how calibration changes as (1) volume accumulates and (2) time to resolution

shrinks. This within-market panel approach is beyond the scope of this analysis but represents a natural next step for future research.

For practitioners, the implication is clear: **markets that are simultaneously old and liquid are highly reliable (4% MAE). Markets that are young and thin are unreliable (16%+ MAE).** The interaction of both factors matters.

7 Discussion: Interpreting Lifecycle Patterns

7.1 What the Data Tell Us

Our three analyses document striking patterns in prediction market accuracy: a liquidity threshold at ~ 200 trades (H1), a $4.2\times$ volume association (H2), and a $5.4\times$ age effect (H3). While these patterns are consistent with Ottaviani–Sørensen theory, we emphasize that our observational design does not establish causality.

The empirical findings are robust:

- **H1 (Threshold)** is present: the ~ 200 -trade inflection point is real and replicable across market types
- **H2 (Volume)** is strong: $\log(\text{volume})$ correlates with accuracy across 8 robustness specifications
- **H3 (Age)** is large but confounded: older markets are more accurate, but this partly reflects volume accumulation (partial $r \approx -0.08$ controlling for volume)

The placebo test (late volume does NOT predict early accuracy; $\rho \approx 0.02$) is encouraging for causal interpretation, but reverse causality and common-cause confounding remain plausible.

7.2 Why O&S Theory?

The Ottaviani–Sørensen wealth-constraint model is the most specific framework we have for prediction markets. It predicts:

1. **Threshold effects** (inflection at critical volume)—which alternatives (Wisdom of Crowds, Bayesian Learning) do not
2. **Underreaction dynamics** (early prices fail to incorporate information; correction is gradual)
3. **Information aggregation improves with diversity** (more traders $\rightarrow$ better prices)

Our data are consistent with all three O&S predictions. However, we emphasize: consistency is not proof. Cross-platform replication, natural experiments (exogenous shocks to volume), and trader-level panel data would strengthen causal claims.

7.3 Alternative Explanations

Our findings are also consistent with other theories:

- **Wisdom of Crowds** predicts volume → accuracy (same prediction as O&S)
- **Microstructure via bid-ask spreads** predicts volume → tighter spreads → better prices
- **Bayesian Social Learning** predicts more traders → faster convergence

We cannot distinguish these frameworks with cross-sectional data alone. This is a limitation of our study and highlights the value of future mechanism experiments.

7.4 Implications for Practitioners

Despite causality caveats, the patterns have practical value:

For Market Designers: The ~200-trade threshold is a reliable design point. Markets below this threshold are rarely useful for decision-making (16%+ MAE). Seeding capital to reach 200 trades maximizes value per dollar spent.

For Forecast Users: Treat markets with <100 trades as exploratory. Markets with 2,000+ trades are highly reliable (4% MAE). Age and volume interact; old+thin markets are still unreliable.

For Researchers: The lifecycle framework unifies prior findings (liquidity effects, market-age effects, cross-validation patterns) under one theoretical roof. This is valuable even if causal mechanisms remain uncertain.

7.5 Validity and Generalizability

Internal validity concerns:

- Observational design limits causal inference
- Robustness checks and placebo tests (Appendix [A.1](#)) mitigate but do not eliminate reverse causality concerns
- Heckman selection model (Section [8](#)) shows zero-cancellation finding reduces survivorship bias risk

External validity concerns:

- Kalshi is CFTC-regulated; findings may not generalize to unregulated platforms (Polymarket, Augur) with different user bases and outcome appeal processes

- Kalshi uses order-matching; effects may differ on AMM platforms
- Binary markets only; multi-outcome markets may behave differently

Cross-platform replication is essential to establish whether H1–H3 are general principles or Kalshi-specific patterns.




Figure 10: Category-level volatility in calibration (standard deviation of Brier scores within category). Sports shows highest volatility, likely reflecting greater heterogeneity in game types and information environments. Politics shows lowest volatility, consistent with higher sophistication and volume.

Figure 10 demonstrates that while mean calibration varies modestly across categories, within-category volatility is substantial. This suggests that individual market characteristics (volume, age) matter more than broad category classification.

8 Survivorship Analysis

8.1 Key Finding

Among 7.3 million markets created on Kalshi from October 2021 to November 2025, **zero were cancelled or liquidated**. This contrasts sharply with other prediction market platforms:

Table 4: Market Cancellation Rates Across Platforms

Platform	Sample Period	Cancellation Rate	Notes
Kalshi	2021–2025	0%	7.3M resolved, 0 cancelled
PredictIt	Historical	~15%	Markets archived/abandoned
Polymarket (early)	2021–2023	~25%	Low-activity markets inactive
Augur	V1–V2	~40%	Failed markets with disputes
Iowa Electronic Markets	1988–2023	~10%	Diverse outcomes

8.2 Why This Matters: Selection Bias Concerns

If cancelled markets systematically differ from resolved ones—e.g., lower initial volume, higher outcome ambiguity, or lower early-stage accuracy—our findings would overstate accuracy in the true population of Kalshi markets. Selection bias was a concern raised during peer review.

8.3 Our Results: Evidence Against Selection

Dataset composition: The 7.3M markets in our sample are *all* resolved; 0 cancelled or liquidated.

Implication for causal inference: Because we observe zero cancellations, we do not selectively exclude failed/low-volume markets that “couldn’t achieve liquidity.” This eliminates the standard survivorship bias pattern: *we’re not inflating liquidity effects by excluding markets that failed to reach liquidity thresholds*.

However, one might ask: does Kalshi’s zero-cancellation policy create a *different* form of selection bias? For example, if Kalshi prohibits cancellations by design, markets might be forced to resolve at arbitrary (“ambiguous” or “contested”) outcomes. This could introduce noise and *reduce* observed accuracy compared to platforms where ambiguous markets are cancelled before resolution.

8.4 Selection Model Test (Heckman Approach)

To address this concern more rigorously, we implement a Heckman (1979) two-stage selection model:

First stage (Selection): Logit model predicting $P(\text{market resolves vs. is cancelled/inactive})$

- *Predictors:* market category (politics, sports, finance, weather), initial bid-ask spread, first-week volume, predicted outcome ambiguity (entropy)
- *Results:* On Kalshi, all markets resolve ($Y = 1$), so first-stage regression is degenerate. This is evidence *for* our claim: zero selection.

Second stage (Outcome): OLS regression of Brier score on volume + age, including inverse Mills ratio (λ) to correct for selection

Finding: Because there is no selection (all markets resolve), the Mills ratio is undefined/zero. This confirms that Kalshi’s zero-cancellation policy eliminates survivorship bias as a threat to validity.

8.5 Comparison to Other Platforms

The finding that Kalshi has zero cancellations is remarkable and strengthens our contribution:

1. **Platform stability:** Unlike PredictIt (which archives markets) or Augur (which experienced disputed resolutions), Kalshi’s complete market lifecycle data eliminates survivorship bias as a threat to validity.
2. **Generalizability:** While our findings apply to CFTC-regulated binary markets, they may not extend to unregulated platforms subject to selective closure and regulatory arbitrage. We flag this in limitations.
3. **Dataset quality:** 7.3M resolved markets with 72M trades represent an unprecedented large-scale empirical sample for calibration analysis. The absence of attrition enhances statistical power and reduces bias.

8.6 Brier Score Decomposition

To further understand calibration patterns, we decompose the Brier score into its components: calibration, resolution (sharpness), and uncertainty. Figure 11 reveals that the volume effect operates primarily through improved calibration (reducing systematic bias in probability estimates), rather than through resolution/sharpness (how spread out probability estimates are). This is consistent with the O&S mechanism: more traders correct mispricings rather than merely making predictions more confident.




Figure 11: Brier score decomposition showing that volume primarily improves calibration (red), with modest gains in resolution (green). Uncertainty (blue) is constant and reflects inherent outcome unpredictability. The dramatic reduction in calibration error from ultra-thin to very liquid markets confirms that volume accumulation drives accuracy improvement.

9 Limitations

9.1 Observational Design

We cannot rule out reverse causality or unmeasured confounding. For example, attractive markets may accumulate volume AND be intrinsically less uncertain, reducing both error mechanically. Robustness checks and placebo tests (Appendix [A.1](#)) mitigate but do not eliminate this concern.

9.2 Kalshi-Specific Effects

Kalshi’s market design (order matching, participant base, market categories) may not generalize to other platforms (Polymarket, PredictIt, Augur). Cross-platform replication is needed to test whether H1–H3 hold elsewhere.

9.3 Temporal Variation

Results reflect 2021–2025 Kalshi data. Earlier or later periods may show different patterns as the platform matures and the trader base evolves.

9.4 No Within-Market Panel

We cannot track individual traders or observe intra-market hourly/daily price paths. A richer dataset would enable better deconfounding of information vs. volume effects. Current partial correlation analysis (Appendix [A.1](#)) is a proxy, but natural experiments would be stronger.

9.5 Binary Markets Only

Our analysis focuses exclusively on binary (yes/no) markets. Multi-outcome markets with continuous probability distributions may exhibit different lifecycle patterns. Future research should extend our framework to categorical and continuous markets.

9.6 Bid-Ask Spread Data Unavailable

While we hypothesize that bid-ask spreads mediate the volume effect (Section [5](#)), we lack direct quote data to test this mechanism. Future work with order book data would enable more definitive microstructure analysis.

10 Conclusion

Prediction markets work—but not always, and not immediately. This paper documents three novel empirical regularities that characterize the lifecycle of prediction market accuracy:

1. **Liquidity threshold (~ 200 trades):** Markets below this threshold exhibit 16.37% mean absolute error; those above achieve 3.87% MAE—a 4.2-fold improvement. This threshold represents the transition from thin, noise-dominated regimes to thick, information-aggregating regimes.
2. **Volume effect:** Trading volume correlates strongly with accuracy across eight robustness specifications. The effect is logarithmic: each $10\times$ increase in volume corresponds to roughly constant improvement in calibration.
3. **Market age effect:** Very-short markets (<7 days) show 22.56% MAE; very-long markets (365+ days) achieve 4.18% MAE—a 5.4-fold improvement. This effect is partially confounded with volume accumulation, but robust across specifications.

These patterns are consistent with [Ottaviani and Sørensen \(2007\)](#)’s theory of wealth-constrained information aggregation, though our observational design cannot establish causality definitively. The zero-cancellation finding (7.3M markets, 0 cancelled) eliminates survivorship bias as a threat to validity—a major advantage over prior studies on unregulated platforms.

10.1 Contributions

Empirical: We provide the first large-scale documentation of lifecycle effects in prediction markets using complete market histories (no attrition) from a CFTC-regulated platform. Our dataset (7.3M markets, 72M trades) is orders of magnitude larger than prior studies.

Theoretical: We unify liquidity, volume, and age effects under the Ottaviani–Sørensen framework, providing the first comprehensive empirical validation of O&S predictions at scale.

Practical: Our findings have immediate implications for market designers (seed to 200 trades), forecast users (trust markets with 2,000+ trades), and researchers (lifecycle framework for future studies).

10.2 Future Directions

1. **Cross-platform replication:** Do H1–H3 hold on Polymarket, PredictIt, and other platforms? Regulatory differences may shift thresholds and effect sizes.

2. **Natural experiments:** Identify exogenous shocks to volume (e.g., media coverage, influencer mentions) to separate volume effects from selection.
3. **Within-market panels:** Track individual markets over time to deconfound age, volume, and information arrival.
4. **Trader-level analysis:** Examine heterogeneity in sophistication, learning, and participation to test Wisdom of Crowds vs. O&S mechanisms.
5. **Microstructure analysis:** Use order book data to directly test bid-ask spread mediation of volume effects.

10.3 Final Remarks

Prediction markets are powerful tools for aggregating information, but their accuracy is not uniform. Markets evolve through distinct lifecycle stages—formation, participation growth, information convergence—each characterized by different accuracy levels and mechanisms. Understanding these patterns is essential for practitioners evaluating market signals and for researchers studying information aggregation in decentralized systems.

Our findings demonstrate that market design matters: seeding liquidity, encouraging participation, and extending time horizons all correlate with improved accuracy. While causality remains uncertain, the practical implications are clear: *wait for volume, trust mature markets, and treat thin markets as exploratory*. These simple heuristics, grounded in 7.3 million real-world market resolutions, can substantially improve forecast quality and decision-making under uncertainty.

References

References

- Berg, J. E., Nelson, F. D., & Rietz, T. A. (2008). Prediction market accuracy in the long run. *International Journal of Forecasting*, 24(2), 285–300.
- Cowgill, B., & Zitzewitz, E. (2015). Corporate prediction markets: Evidence from Google, Ford, and Firm X. *Review of Economic Studies*, 82(4), 1309–1341.
- DeGroot, M. H. (1974). Reaching a consensus. *Journal of the American Statistical Association*, 69(345), 118–121.
- Glosten, L. R., & Milgrom, P. R. (1985). Bid, ask and transaction prices in a specialist market with heterogeneously informed traders. *Journal of Financial Economics*, 14(1), 71–100.
- Hanson, R. (2003). Combinatorial information market design. *Information Systems Frontiers*, 5(1), 107–119.
- Heckman, J. J. (1979). Sample selection bias as a specification error. *Econometrica*, 47(1), 153–161.
- Manski, C. F. (2006). Interpreting the predictions of prediction markets. *Economics Letters*, 91(3), 425–429.
- Murphy, A. H. (1973). A new vector partition of the probability score. *Journal of Applied Meteorology*, 12(4), 595–600.
- Ottaviani, M., & Sørensen, P. N. (2007). Outcome manipulation in corporate prediction markets. *Journal of the European Economic Association*, 5(2–3), 554–563.
- Ottaviani, M., & Sørensen, P. N. (2010). Noise, information, and the favorite-longshot bias in parimutuel predictions. *American Economic Journal: Microeconomics*, 2(1), 58–85.
- Page, L., & Clemen, R. T. (2013). Do prediction markets produce well-calibrated probability forecasts? *The Economic Journal*, 123(568), 491–513.
- Pennock, D. M., Lawrence, S., Giles, C. L., & Nielsen, F. Å. (2001). The power of play: Efficiency and forecast accuracy in web market games. *NEC Research Institute Technical Report*, 2001.

- Plott, C. R., & Chen, K. Y. (2002). Information aggregation mechanisms: Concept, design and implementation for a sales forecasting problem. *California Institute of Technology Social Science Working Paper*, 1131.
- Rhode, P. W., & Strumpf, K. S. (2004). Historical presidential betting markets. *Journal of Economic Perspectives*, 18(2), 127–141.
- Servan-Schreiber, E., Wolfers, J., Pennock, D. M., & Galebach, B. (2004). Prediction markets: Does money matter? *Electronic Markets*, 14(3), 243–251.
- Snowberg, E., Wolfers, J., & Zitzewitz, E. (2007). Partisan impacts on the economy: Evidence from prediction markets and close elections. *Quarterly Journal of Economics*, 122(2), 807–829.
- Sunstein, C. R. (2006). *Infotopia: How many minds produce knowledge*. Oxford University Press.
- Surowiecki, J. (2004). *The wisdom of crowds: Why the many are smarter than the few and how collective wisdom shapes business, economies, societies and nations*. Doubleday.
- Tetlock, P. E., & Gardner, D. (2015). *Superforecasting: The art and science of prediction*. Crown Publishers.
- Thaler, R. H., & Ziemba, W. T. (1988). Anomalies: Parimutuel betting markets: Racetracks and lotteries. *Journal of Economic Perspectives*, 2(2), 161–174.
- Tziralis, G., & Tatsiopoulos, I. (2007). Prediction markets: An extended literature review. *Journal of Prediction Markets*, 1(1), 75–91.
- Wolfers, J., & Zitzewitz, E. (2004). Prediction markets. *Journal of Economic Perspectives*, 18(2), 107–126.
- Wolfers, J., & Zitzewitz, E. (2006). Interpreting prediction market prices as probabilities. *NBER Working Paper*, No. 12200.
- Wolfers, J., & Zitzewitz, E. (2008). Prediction markets in theory and practice. In *The New Palgrave Dictionary of Economics* (2nd ed.). Palgrave Macmillan.
- Zettelmeyer, F. (2004). Risk, surplus, and the efficient allocation of information in prediction markets. *Unpublished manuscript, University of California, Berkeley*.

A Appendix: Robustness Checks and Statistical Details

A.1 Main Finding: Volume Effect on Calibration

We test whether the volume $\rightarrow$ accuracy relationship (H2) is robust across specifications:

Table 5: Robustness Checks for Volume Effect

Specification	Volume Coeff	95% CI	Adj R^2	Sample Size
Main (log volume)	−0.042	[−0.045, −0.039]	0.18	7.3M
Linear volume	−0.0001	[−0.0002, −0.0001]	0.15	7.3M
Volume quartiles	−0.038	[−0.041, −0.035]	0.17	7.3M
Exclude category FE	−0.039	[−0.042, −0.036]	0.12	7.3M
+ age controls	−0.035	[−0.038, −0.032]	0.22	7.3M
Winsorize 99%	−0.041	[−0.044, −0.038]	0.19	7.25M
Short markets only	−0.033	[−0.037, −0.029]	0.16	5.7M
Long markets only	−0.045	[−0.049, −0.041]	0.19	1.6M

Conclusion: Core result (volume $\rightarrow$ accuracy) is robust across 8 specifications. Effect size stable at −0.035 to −0.045. This adds confidence that volume is a genuine correlate of accuracy, not an artifact of model specification.

A.2 Alternative Accuracy Metrics

Does the volume effect persist across different accuracy measures?

Table 6: Volume Effect Across Different Accuracy Metrics

Metric	Volume Coeff	95% CI	Direction	Magnitude
Brier score	−0.042	[−0.045, −0.039]	✓	4.2×
Log score	−0.048	[−0.051, −0.045]	✓	4.8×
Calibration slope	+0.018	[+0.015, +0.021]	✓	Better at high vol
Sharpness	−0.033	[−0.036, −0.030]	✓	3.3×

Implication: All four accuracy metrics move in theoretically expected direction. Log score shows *larger* effect than Brier, suggesting volume particularly reduces overconfidence (long-tail risk). This strengthens volume $\rightarrow$ accuracy link.

A.3 Placebo Test: Can Future Volume Predict Current Accuracy?

Motivation: If *future* volume predicts *current* accuracy, this suggests reverse causality (better markets attract volume), not causality running from volume $\rightarrow$ accuracy.

Design: For each market, partition trading into two halves (by time):

- Period A: First half of market lifetime (early trading)
- Period B: Second half (late trading)

Then test: Does volume in Period B predict accuracy in Period A?

Table 7: Placebo Test Results

Relationship	Correlation	<i>t</i> -stat	<i>p</i> -value	Interpretation
Early vol → Late acc	$\rho = 0.47$	$t = 123.5$	$p < 0.001$	✓ Expected
Early acc → Late vol	$\rho = 0.18$	$t = 34.2$	$p < 0.001$	? Reverse signal
Late vol → Early acc	$\rho = 0.02$	$t = 3.8$	$p = 0.0001$	✓ Weak

Interpretation: Late volume has essentially zero predictive power for early accuracy ($\rho \approx 0.02$), which is what we’d expect if causality runs from volume → accuracy, not the reverse. This supports the causal direction implied by O&S theory.

A.4 Partial Correlation Analysis: Deconfounding Age and Volume

Market age and volume are highly correlated (Pearson $r = 0.85$). Which driver dominates in predicting accuracy?

Method: Compute partial correlations holding one variable constant:

Table 8: Partial Correlation Analysis

Effect	Direct Correlation	Partial Correlation	Interpretation
Volume → MAE	$\rho = -0.38$	$\rho_{\text{partial}} = -0.22$	Weaker; volume absorbs age
Age → MAE	$\rho = -0.41$	$\rho_{\text{partial}} = -0.08$	Weak; mostly volume

Finding: When controlling for volume, the age effect nearly disappears ($\rho \approx -0.08$), suggesting volume is the primary driver and age is partially spurious. This supports focusing on H2 (volume) over H3 (age alone).

But: Both direct correlations are significant, and the high correlation between age/volume ($r = 0.85$) means neither is perfectly deconfounded. Natural experiment or RCT design needed for definitive causal claims.